

```
In [159]: from google.colab import drive
drive.mount('/content/drive')

import sys
sys.path.append('drive/MyDrive/USC Spring 2024/Machine Learning/Project/')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
In [156]: import numpy as np
import pandas as pd
import pickle
import matplotlib.pyplot as plt
```

```
In [157]: from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold
from sklearn.linear_model import Perceptron
from sklearn.neighbors import NeighborhoodComponentsAnalysis, KNeighborsClassifier
from sklearn.metrics import f1_score, ConfusionMatrixDisplay
```

```
In [4]: def load_data(filename):
        #Loading the file into data
        data = np.array(pd.read_csv(filename))
        return data
```

```
In [5]: drybeans_dataset_train_filename = 'drive/MyDrive/USC Spring 2024/Machine Learning/Project/dry_bean_classification_train.csv'
drybeans_dataset_test_filename = 'drive/MyDrive/USC Spring 2024/Machine Learning/Project/dry_bean_classification_test.csv'
```

```
In [6]: drybeans_dataset_train = load_data(drybeans_dataset_train_filename)
#Splitting into X and y dataset
drybeans_dataset_train_X = drybeans_dataset_train[:, :-1]
drybeans_dataset_train_y = drybeans_dataset_train[:, -1]

drybeans_dataset_test = load_data(drybeans_dataset_test_filename)
drybeans_dataset_test_X = drybeans_dataset_test[:, :-1]
drybeans_dataset_test_y = drybeans_dataset_test[:, -1]
```

```
In [8]: def standardize_training_dataset(dataset):
        dataset_copy = dataset.copy()
        means = []
        std_devs = []

        for i in range(dataset.shape[1]):
            column = dataset[:, i]
            mean = np.mean(column)
            std_dev = np.std(column)

            means.append(mean)
            std_devs.append(std_dev)

            dataset_copy[:, i] = (column-mean)/std_dev

        return means, std_devs, dataset_copy
```

```
In [9]: def standardize_test_dataset(dataset, means, std_devs, ):
        dataset_test_copy = dataset.copy()

        for i in range(dataset.shape[1]):
            column = dataset[:, i]
            dataset_test_copy[:, i] = (column-means[i])/std_devs[i]

        return dataset_test_copy
```

```
In [10]: def compute_test_accuracy(predicted, actual):
        return np.sum(predicted == actual)/len(actual)
```

```
In [11]: def compute_F1_macro(actual, predicted):
        return f1_score(y_true=actual, y_pred=predicted, average='macro')
```

```
In [12]: def compute_F1_weighted(actual, predicted):
        return f1_score(y_true=actual, y_pred=predicted, average='weighted')
```

```
In [13]: def create_confusion_matrix(model, dataset, labels):

        disp = ConfusionMatrixDisplay.from_estimator(
            model,
            dataset,
            drybeans_dataset_test_y,
            cmap=plt.cm.Blues,
            normalize=None)

        plt.figure(figsize=(20, 20))

        disp.ax_.set_title('Confusion Matrix')

        plt.show()
```

```
In [16]: training_means, training_std_devs, standardized_drybeans_dataset_train_X = standardize_training_dataset(drybeans_dataset_train_X)
        standardized_drybeans_dataset_test_X = standardize_test_dataset(drybeans_dataset_test_X, training_means, training_std_devs)
```

## 2. Support Vector Classifier (SVC)

```
In [46]: def model_selection_svc(C_range, gamma_range, dataset, labels, k):
         average_validation_classification_accuracy_C = {}

         for C in C_range:
             average_validation_classification_accuracy_C[C] = []

             for gamma in gamma_range:
                 KF = KFold(n_splits=k)
                 gamma_valid_accuracy = []

                 for i, (train, valid) in enumerate(KF.split(dataset, labels)):
                     svc_model = SVC(C=C, gamma=gamma)
                     svc_model.fit(dataset[train], labels[train])
                     valid_accuracy = svc_model.score(dataset[valid], labels[valid])

                     gamma_valid_accuracy.append(valid_accuracy)

                 average_validation_classification_accuracy_C[C].append(
                     np.mean(gamma_valid_accuracy))

         return average_validation_classification_accuracy_C
```

```
In [147]: def find_best_parameter_svc(C_range, gamma_range, average_validation_classification_accuracy):
         best_gamma = 0
         best_C = 0
         best_avg_valid_accuracy = float('-inf')

         for C in C_range:
             for i in range(len(average_validation_classification_accuracy[C])):
                 print('C =', C, ', Gamma =', gamma_range[i], ', Average Validation Accuracy =', average_validation_classification_accuracy[C][i])
                 if average_validation_classification_accuracy[C][i] >= best_avg_valid_accuracy:
                     best_avg_valid_accuracy = average_validation_classification_accuracy[C][i]
                     best_C = C
                     best_gamma = gamma_range[i]

         return best_C, best_avg_valid_accuracy, best_gamma
```

```
In [48]: def find_best_validation_accuracy(mean_validation_accuracy, C_vals):
        best_mean_validation_accuracy_C = []

        for C in C_vals:
            best_mean_validation_accuracy_C.append(max(mean_validation_
accuracy[C]))
            best_index = np.argmin(mean_validation_accuracy[C])

        return best_mean_validation_accuracy_C
```

```
In [49]: def plot_ave_valid_accuracy_vs_C(C_range, validation_accuracy):
        plt.plot(C_range, validation_accuracy)
        plt.title('Average Validation Accuracy vs C')
        plt.xlabel('C')
        plt.xscale('log')
        plt.ylabel('Average Validation Accuracy')
        plt.show()
```

```
In [50]: def plot_ave_valid_accuracy_vs_gamma(gamma_range, validation_accu
cy):
        plt.plot(gamma_range, validation_accuracy)
        plt.title('Average Validation Accuracy vs Gamma')
        plt.xlabel('Gamma')
        plt.xscale('log')
        plt.ylabel('Average Validation Accuracy')
        plt.show()
```

```
In [51]: C_range = [0.1, 1, 10, 100]
        gamma_range = [0.1, 1, 10, 100]
        K = 5



        average_validation_classification_accuracy_C = model_selection_svc
(C_range, gamma_range, standardized_drybeans_dataset_train_X, drybe
ans_dataset_train_y, K)
```

```
In [148]: best_C, best_avg_valid_accuracy, best_gamma = find_best_parameter_s
vc(C_range, gamma_range, average_validation_classification_accuracy
_C)

print('Best gamma:', best_gamma)
print('Best C:', best_C)
print('Best Average Validation Accuracy:', best_avg_valid_accuracy)

C = 0.1 , Gamma = 0.1 , Average Validation Accuracy = 0.92432758705
7783
C = 0.1 , Gamma = 1 , Average Validation Accuracy = 0.8798788823635
36
C = 0.1 , Gamma = 10 , Average Validation Accuracy = 0.285518525126
8268
C = 0.1 , Gamma = 100 , Average Validation Accuracy = 0.25851891782
906106
C = 1 , Gamma = 0.1 , Average Validation Accuracy = 0.9284602402696
528
C = 1 , Gamma = 1 , Average Validation Accuracy = 0.919736092287977
6
C = 1 , Gamma = 10 , Average Validation Accuracy = 0.65460670090895
18
C = 1 , Gamma = 100 , Average Validation Accuracy = 0.2669679211626
011
C = 10 , Gamma = 0.1 , Average Validation Accuracy = 0.930847962651
5288
C = 10 , Gamma = 1 , Average Validation Accuracy = 0.91450062490693
87
C = 10 , Gamma = 10 , Average Validation Accuracy = 0.6739833293472
579
C = 10 , Gamma = 100 , Average Validation Accuracy = 0.268804637176
45825
C = 100 , Gamma = 0.1 , Average Validation Accuracy = 0.92882746536
64891
C = 100 , Gamma = 1 , Average Validation Accuracy = 0.8996229889828
253
C = 100 , Gamma = 10 , Average Validation Accuracy = 0.673983329347
258
C = 100 , Gamma = 100 , Average Validation Accuracy = 0.26880463717
645825
Best gamma: 0.1
Best C: 10
Best Average Validation Accuracy: 0.9308479626515288
```

```
In [53]: best_mean_validation_accuracy_C = find_best_validation_accuracy(average_validation_classification_accuracy_C, C_range)
plot_ave_valid_accuracy_vs_C(C_range, best_mean_validation_accuracy_C)
plot_ave_valid_accuracy_vs_gamma(gamma_range, average_validation_classification_accuracy_C[best_C])
```



```
In [54]: SVC_model_final = SVC(C=best_C, gamma=best_gamma)
SVC_model_final.fit(standardized_drybeans_dataset_train_X, drybeans_dataset_train_y)
```

```
Out[54]: SVC
SVC(C=10, gamma=0.1)
```

```
In [55]: predictions = SVC_model_final.predict(standardized_drybeans_dataset_test_X)
test_accuracy_final_model = compute_test_accuracy(predictions, drybeans_dataset_test_y)
print('Classification Test Accuracy:', test_accuracy_final_model*100, '%')
```

Classification Test Accuracy: 93.27700220426158 %

```
In [56]: F1_macro_SVC = compute_F1_macro(drybeans_dataset_test_y, predictions)
F1_weighted_SVC = compute_F1_weighted(drybeans_dataset_test_y, predictions)


print('Macro-Averaged F1 Score =', F1_macro_SVC)
print('Weighted-Averaged F1 Score =', F1_weighted_SVC)
```

Macro-Averaged F1 Score = 0.945424463981152

Weighted-Averaged F1 Score = 0.9329631451281314



```
In [57]: create_confusion_matrix(SVC_model_final, standardized_drybeans_data
set_test_X, drybeans_dataset_test_y)
```



<Figure size 2000x2000 with 0 Axes>

### 3. Random Forest Classifier

```
In [149]: def model_selection_rfc(max_depth_range, tree_range, dataset, labels, k):
    average_validation_classification_accuracy_tree = {}

    for tree in tree_range:
        average_validation_classification_accuracy_tree[tree] = []

        for max_depth in max_depth_range:
            KF = KFold(n_splits=k)
            depth_valid_accuracy = []

            for i, (train, valid) in enumerate(KF.split(dataset, labels)):
                rfc_model = RandomForestClassifier(n_estimators=tree, max_depth=max_depth)
                rfc_model.fit(dataset[train], labels[train])
                valid_accuracy = rfc_model.score(dataset[valid], labels[valid])

                depth_valid_accuracy.append(valid_accuracy)

            average_validation_classification_accuracy_tree[tree].append(np.mean(depth_valid_accuracy))

    return average_validation_classification_accuracy_tree
```

```
In [163]: def find_best_parameter_rf(max_depth_range, tree_range, average_validation_classification_accuracy):
    best_max_depth = 0
    best_tree = 0
    best_avg_valid_accuracy = float('-inf')

    for tree in tree_range:
        for i in range(len(average_validation_classification_accuracy[tree])):
            print('Tree =', tree, ', Max Depth =', max_depth_range[i], ', Average Validation Accuracy =', average_validation_classification_accuracy[tree][i])
            if average_validation_classification_accuracy[tree][i] >= best_avg_valid_accuracy:
                best_avg_valid_accuracy = average_validation_classification_accuracy[tree][i]
                best_tree = tree
                best_max_depth = max_depth_range[i]

    return best_tree, best_avg_valid_accuracy, best_max_depth
```

```
In [151]: def find_best_validation_accuracy(mean_validation_accuracy, tree_vals):
            best_mean_validation_accuracy_tree = []

            for tree in tree_vals:
                best_mean_validation_accuracy_tree.append(max(mean_validation_accuracy[tree]))
                best_index = np.argmin(mean_validation_accuracy[tree])

            return best_mean_validation_accuracy_tree
```

```
In [152]: def plot_ave_valid_accuracy_vs_tree(tree_range, validation_accuracy):
            plt.plot(tree_range, validation_accuracy)
            plt.title('Average Validation Accuracy vs Number of Trees')
            plt.xlabel('Number of Trees')
            plt.xscale('log')
            plt.ylabel('Average Validation Accuracy')
            plt.show()
```

```
In [153]: def plot_ave_valid_accuracy_vs_max_depth(max_depth_range, validation_accuracy):
            plt.plot(max_depth_range, validation_accuracy)
            plt.title('Average Validation Accuracy vs Max Depth')
            plt.xlabel('Max Depth')
            plt.xscale('log')
            plt.ylabel('Average Validation Accuracy')
            plt.show()
```

```
In [161]: tree_range = [1, 10, 100]
            max_depth_range = [1, 10, 100]
            K = 5



            average_validation_classification_accuracy_tree = model_selection_rfc(max_depth_range, tree_range, standardized_drybeans_dataset_train_X, drybeans_dataset_train_y, K)
```

```
In [164]: best_tree, best_avg_valid_accuracy, best_max_depth = find_best_parameter_rf(max_depth_range, tree_range, average_validation_classification_accuracy_tree)

print('Best Number of Trees:', best_tree)
print('Best Max Depth:', best_max_depth)
print('Best Average Validation Accuracy', best_avg_valid_accuracy)

Tree = 1 , Max Depth = 1 , Average Validation Accuracy = 0.4055469717849139
Tree = 1 , Max Depth = 10 , Average Validation Accuracy = 0.894757066636634
Tree = 1 , Max Depth = 100 , Average Validation Accuracy = 0.8841031098558139
Tree = 10 , Max Depth = 1 , Average Validation Accuracy = 0.4992138362790219
Tree = 10 , Max Depth = 10 , Average Validation Accuracy = 0.921113418394915
Tree = 10 , Max Depth = 100 , Average Validation Accuracy = 0.9156035656181812
Tree = 100 , Max Depth = 1 , Average Validation Accuracy = 0.437136302263458
Tree = 100 , Max Depth = 10 , Average Validation Accuracy = 0.9220317764018435
Tree = 100 , Max Depth = 100 , Average Validation Accuracy = 0.9235932634061836
Best Number of Trees: 100
Best Max Depth: 100
Best Average Validation Accuracy 0.9235932634061836
```

```
In [165]: best_mean_validation_accuracy_tree = find_best_validation_accuracy  
(average_validation_classification_accuracy_tree, tree_range)  
plot_ave_valid_accuracy_vs_tree(tree_range, best_mean_validation_ac  
curacy_tree)  
plot_ave_valid_accuracy_vs_max_depth(max_depth_range, average_vali  
dation_classification_accuracy_tree[best_tree])
```



```
In [166]: rf_model_final = RandomForestClassifier(n_estimators=best_tree, max_
          _depth=best_max_depth)
          rf_model_final.fit(standardized_drybeans_dataset_train_X, drybeans_
          dataset_train_y)
```

```
Out[166]: ▼      RandomForestClassifier
          RandomForestClassifier(max_depth=100)
```

```
In [167]: predictions = rf_model_final.predict(standardized_drybeans_dataset_
          test_X)
          test_accuracy_rf = compute_test_accuracy(drybeans_dataset_test_y, p
          redictions)
          print('Classification Test Accuracy:', test_accuracy_rf*100, '%')
```

Classification Test Accuracy: 91.9911829537105 %


```
In [168]: F1_macro_RF = compute_F1_macro(drybeans_dataset_test_y, prediction
          s)
          F1_weighted_RF = compute_F1_weighted(drybeans_dataset_test_y, predi
          ctions)

          print('Macro-Averaged F1 Score =', F1_macro_RF)
          print('Weighted-Averaged F1 Score =', F1_weighted_RF)
```

Macro-Averaged F1 Score = 0.9320815939173662

Weighted-Averaged F1 Score = 0.9199509695063591

```
In [169]: create_confusion_matrix(rf_model_final, standardized_drybeans_datas
et_test_X, drybeans_dataset_test_y)
```



<Figure size 2000x2000 with 0 Axes>



## 4. KNN Classifier

```
In [34]: def model_selection_knn(neighbor_range, dataset, labels, k):
    best_neighbor = 0
    average_validation_classification_accuracy = []

    for neighbor in neighbor_range:
        KF = KFold(n_splits=k)
        neighbor_valid_accuracy = []

        for i, (train, valid) in enumerate(KF.split(dataset, labels)):

            knn_model = KNeighborsClassifier(n_neighbors=neighbor)

            knn_model.fit(dataset[train], labels[train])

            valid_accuracy = knn_model.score(dataset[valid], labels[valid])

            neighbor_valid_accuracy.append(valid_accuracy)

            average_validation_classification_accuracy.append(np.mean(neighbor_valid_accuracy))

        for i in range(len(average_validation_classification_accuracy)):
            print('Number of Neighbors =', neighbor_range[i], ', Average Validation RMSE =', average_validation_classification_accuracy[i])

            best_neighbor_index = np.argmax(average_validation_classification_accuracy)
            best_neighbor = neighbor_range[best_neighbor_index]
            best_average_validation_accuracy = average_validation_classification_accuracy[best_neighbor_index]

        return best_neighbor, average_validation_classification_accuracy, best_average_validation_accuracy
```

```
In [15]: def plot_ave_valid_accuracy_vs_neighbors(neighbor_range, validation_accuracy):
    plt.plot(neighbor_range, validation_accuracy)
    plt.title('Average Validation Accuracy vs Number of Neighbors')
    plt.xlabel('Gamma')
    plt.xscale('log')
    plt.ylabel('Average Validation Accuracy')
    plt.show()
```

```
In [35]: neighbor_range = [1, 10, 100, 1000]
K = 5

best_neighbor, average_validation_classification_accuracy, best_ave
rage_validation_accuracy = model_selection_knn(neighbor_range, stan
dardized_drybeans_dataset_train_X, drybeans_dataset_train_y, K)
print('Best number of neighbors:', best_neighbor)
print('Best average validation accuracy:', best_average_validation_
accuracy)
```

Number of Neighbors = 1 , Average Validation RMSE = 0.9034802444624134

Number of Neighbors = 10 , Average Validation RMSE = 0.9251543286036125


Number of Neighbors = 100 , Average Validation RMSE = 0.9176237254576922

Number of Neighbors = 1000 , Average Validation RMSE = 0.8501239057801466

Best number of neighbors: 10

Best average validation accuracy: 0.9251543286036125

```
In [36]: plot_ave_valid_accuracy_vs_neighbors(neighbor_range, average_valida
tion_classification_accuracy)
```



```
In [58]: knn_model_final = KNeighborsClassifier(n_neighbors=best_neighbor)
knn_model_final.fit(standardized_drybeans_dataset_train_X, drybeans_
_dataset_train_y)
```

```
Out [58]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=10)
```

```
In [59]: predictions = knn_model_final.predict(standardized_drybeans_dataset_test_X)
test_accuracy_knn = compute_test_accuracy(drybeans_dataset_test_y, predictions)
print('Classification Test Accuracy:', test_accuracy_knn*100, '%')
```

Classification Test Accuracy: 92.72593681116827 %


```
In [60]: F1_macro_KNN = compute_F1_macro(drybeans_dataset_test_y, predictions)
F1_weighted_KNN = compute_F1_weighted(drybeans_dataset_test_y, predictions)

print('Macro-Averaged F1 Score =', F1_macro_KNN)
print('Weighted-Averaged F1 Score =', F1_weighted_KNN)
```

Macro-Averaged F1 Score = 0.9392816197104452

Weighted-Averaged F1 Score = 0.9272865713896141

```
In [45]: create_confusion_matrix(knn_model_final, standardized_drybeans_dataset_test_X, drybeans_dataset_test_y)
```



<Figure size 2000x2000 with 0 Axes>

## 5. Perceptron

```
In [118]: def model_selection_perceptron(dataset, labels, k, eta_range, alpha_range):
    average_validation_classification_accuracy_alpha = {}

    for alpha in alpha_range:
        average_validation_classification_accuracy_alpha[alpha] = []
        for eta in eta_range:
            KF = KFold(n_splits=k)
            eta_valid_accuracy = []

            for i, (train, valid) in enumerate(KF.split(dataset, labels)):
                perceptron_model = Perceptron(penalty='l2', alpha=alpha, eta0=eta)
                perceptron_model.fit(dataset[train], labels[train])
                valid_accuracy = perceptron_model.score(dataset[valid], labels[valid])

                eta_valid_accuracy.append(valid_accuracy)

            average_validation_classification_accuracy_alpha[alpha].append(np.mean(eta_valid_accuracy))

    return average_validation_classification_accuracy_alpha
```

```
In [145]: def find_best_parameter_perceptron(alpha_range, eta_range, average_validation_classification_accuracy):
    best_alpha = 0
    best_eta = 0
    best_avg_valid_accuracy = float('-inf')

    for alpha in alpha_range:
        for i in range(len(average_validation_classification_accuracy[alpha])):
            print('Alpha =', alpha, ', Eta =', eta_range[i], ', Average Validation Accuracy =', average_validation_classification_accuracy[alpha][i])
            if average_validation_classification_accuracy[alpha][i] >= best_avg_valid_accuracy:
                best_avg_valid_accuracy = average_validation_classification_accuracy[alpha][i]
                best_alpha = alpha
                best_eta = eta_range[i]

    return best_alpha, best_avg_valid_accuracy, best_eta
```

```
In [65]: def find_best_validation_accuracy(mean_validation_accuracy, alpha_vals):
        best_mean_validation_accuracy_alpha = []

        for alpha in alpha_vals:
            best_mean_validation_accuracy_alpha.append(max(mean_validation_accuracy[alpha]))
            best_index = np.argmax(mean_validation_accuracy[alpha])

        return best_mean_validation_accuracy_alpha
```

```
In [66]: def plot_ave_valid_accuracy_vs_alpha(alpha_range, validation_accuracy):
        plt.plot(alpha_range, validation_accuracy)
        plt.title('Average Validation Accuracy vs Alpha')
        plt.xlabel('Alpha')
        plt.xscale('log')
        plt.ylabel('Average Validation Accuracy')
        plt.show()
```

```
In [67]: def plot_ave_valid_accuracy_vs_eta(eta_range, validation_accuracy):
        plt.plot(eta_range, validation_accuracy)
        plt.title('Average Validation Accuracy vs Eta')
        plt.xlabel('Eta')
        plt.xscale('log')
        plt.ylabel('Average Validation Accuracy')
        plt.show()
```

```
In [120]: alpha_range = [0.01, 0.1, 1, 10]
eta_range = [0.01, 0.1, 1, 10]
K = 5



average_validation_classification_accuracy_alpha = model_selection_
perceptron(standardized_drybeans_dataset_train_X, drybeans_dataset_
train_y, K, eta_range, alpha_range)
```

```
In [146]: best_alpha, best_avg_valid_accuracy, best_eta = find_best_parameter
_perceptron(alpha_range, eta_range, average_validation_classification_accuracy_alpha)

print('Best alpha:', best_alpha)
print('Best eta:', best_eta)
print('Best Average Validation Accuracy', best_avg_valid_accuracy)
```

```
Alpha = 0.01 , Eta = 0.01 , Average Validation Accuracy = 0.8703294
691602202
Alpha = 0.01 , Eta = 0.1 , Average Validation Accuracy = 0.80465573
59623714
Alpha = 0.01 , Eta = 1 , Average Validation Accuracy = 0.6042837866
281304
Alpha = 0.01 , Eta = 10 , Average Validation Accuracy = 0.384872359
1196553
Alpha = 0.1 , Eta = 0.01 , Average Validation Accuracy = 0.82110434
95463256
Alpha = 0.1 , Eta = 0.1 , Average Validation Accuracy = 0.614745104
1926342
Alpha = 0.1 , Eta = 1 , Average Validation Accuracy = 0.38891883717
957965
Alpha = 0.1 , Eta = 10 , Average Validation Accuracy = 0.1025852967
3905295
Alpha = 1 , Eta = 0.01 , Average Validation Accuracy = 0.7176968245
953923
Alpha = 1 , Eta = 0.1 , Average Validation Accuracy = 0.38745010551
49988
Alpha = 1 , Eta = 1 , Average Validation Accuracy = 0.1025852967390
5295
Alpha = 1 , Eta = 10 , Average Validation Accuracy = 0.102585296739
05295
Alpha = 10 , Eta = 0.01 , Average Validation Accuracy = 0.503346742
5750384
Alpha = 10 , Eta = 0.1 , Average Validation Accuracy = 0.1259161118
8512677
Alpha = 10 , Eta = 1 , Average Validation Accuracy = 0.102585296739
05295
Alpha = 10 , Eta = 10 , Average Validation Accuracy = 0.10258529673
905295
Best alpha: 0.01
Best eta: 0.01
Best Average Validation Accuracy 0.8703294691602202
```

```
In [122]: best_mean_validation_accuracy_alpha = find_best_validation_accuracy  
(average_validation_classification_accuracy_alpha, alpha_range)  
plot_ave_valid_accuracy_vs_alpha(alpha_range, best_mean_validation_  
accuracy_alpha)  
plot_ave_valid_accuracy_vs_eta(eta_range, average_validation_classi  
fication_accuracy_alpha[best_alpha])
```





```
In [131]: perceptron_model_final = Perceptron(alpha=best_alpha, eta0=best_eta)
perceptron_model_final.fit(standardized_drybeans_dataset_train_X, drybeans_dataset_train_y)
```

```
Out[131]:
```

▼

Perceptron

Perceptron(alpha=0.01, eta0=0.01)

```
In [132]: predictions = perceptron_model_final.predict(standardized_drybeans_dataset_test_X)
test_accuracy_perceptron = compute_test_accuracy(drybeans_dataset_test_y, predictions)
print('Classification Test Accuracy:', test_accuracy_perceptron*100, '%')
```

Classification Test Accuracy: 88.94195444526083 %


```
In [133]: F1_macro_perceptron = compute_F1_macro(drybeans_dataset_test_y, predictions)
F1_weighted_perceptron = compute_F1_weighted(drybeans_dataset_test_y, predictions)

print('Macro-Averaged F1 Score =', F1_macro_perceptron)
print('Weighted-Averaged F1 Score =', F1_weighted_perceptron)
```

Macro-Averaged F1 Score = 0.9033947812491739

Weighted-Averaged F1 Score = 0.8889073871701505

```
In [134]: create_confusion_matrix(perceptron_model_final, standardized_drybeans_dataset_test_X, drybeans_dataset_test_y)
```



<Figure size 2000x2000 with 0 Axes>